



create table S (  
id serial,  
t integer,  
primary key (id),  
foreign key (t) references T(id)  
);

create table T (  
id serial,  
c text not null,  
primary key (id)  
);



create table Uses (  
recipe integer,  
ingredient integer,  
amount integer not null,  
unit text not null,  
primary key (recipe, ingredient),  
foreign key (recipe) references Recipes(id),  
foreign key (ingredient) references Ingredients(id),  
constraint positive\_amount check (amount > 0)  
);

confusing




ER-style  
Person: ID, name  
Employee: ID, salary  
Manager: ID, bonus

O-O style  
Person: ID, name  
Employee: ID, name, salary  
Manager: ID, name, salary, bonus

Single Table  
Person: ID, name, salary, bonus  
(class depends on null values)

ER-style  
create table A (  
id integer,  
x text,  
primary key (id)  
);  
create table B (  
a integer references A(id),  
y text,  
primary key (a)  
);  
create table C (  
a integer references A(id),  
primary key (a)  
);

Single Table  
create table A (  
id integer,  
x text,  
y text,  
b boolean, -- true if A is a B  
c boolean, -- true if A is a C  
primary key (x)  
);

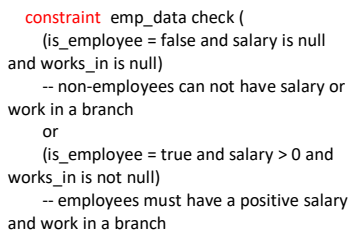


ER-style  
create table People (  
id integer primary key,  
name text not null,  
lives\_in text not null  
);


create table Customers (  
id integer primary key references People(id)  
);

create table Employees (  
id integer primary key references People(id),  
salary integer not null check (salary > 0),  
works\_in integer not null references Branches(id)  
);


Single Table  
create table People (  
id integer primary key,  
name text not null,  
lives\_in text not null,  
  
is\_customer boolean not null,  
is\_employee boolean not null,  
salary integer,  
works\_in integer references branches(id),  
  
constraint emp\_data check (  
(is\_employee = false and salary is null and works\_in is null)  
-- non-employees can not have salary or work in a branch  
or  
(is\_employee = true and salary > 0 and works\_in is not null)  
-- employees must have a positive salary and work in a branch  
),  
constraint total\_participation check (  
is\_customer = true or is\_employee = true  
)  
);



overlapping



Single Table  
create table R (  
id integer primary key,  
is\_a char(1) not null check (is\_a in ('A','B','C')),  
x text,  
y text  
);



Superkey -> Candidate key -> Primary key  
Functional Dependency example  
-> (Position -> Salary)  
Functional Dependencies

hold O: A->B  
hold X: A->C R->C AR->C  
  
Closure: Largest collection of dependencies that can be derived from set F (aka F+)  
F = {A->B, BC->F, BD->EG, AD->C, BEG->FA}  
ACEG+ = {A, B, C, E, F, G}

Minimal Cover: Smallest set of FDs which covers entire FD set

F = {A->BC, B->C, A->B, AB->C}  
1) Convert FDs into canonical form  
{A->B, A->C, B->C, AB->C}  
2) For multi-attribute LHS, remove redundant attributes  
{A->B, A->C, B->C, AB->C}  
3) Eliminate redundant FDs  
{A->B, A->C, B->C}

For all FDs X->Y...  
1) X->Y is trivial (is Y a subset of X?)  
2) X is a superkey  
3) Y is single attribute & part of candidate key  
1,2: BCNF, 1,2,3: 3NF

F = {AB->C, AB->D, C->A, D->B}  
AB+ = ABCD, C+ = AC, CD+ = ABCD

	3NF?	BCNF?
AB->C	2	2
AB->D	2	2
C->A	3	X
D->B	3	X

Transaction: 'unit of work', may change multiple databases at once.

READ, WRITE, BEGIN, COMMIT, ABORT  
Syntax: R(X), W(Y), etc

No concurrency  
T1: R(X) W(X) R(Y) W(Y)  
T2: R(X) W(X) R(Y) W(Y)  
With concurrency  
T1: R(X) W(X) R(Y) W(Y)  
T2: R(X) W(X) R(Y) W(Y)

Two concurrent transfers from same source account:  
• T1 transfers \$200 X->Y, T2 transfers \$100 X->Y  
• initial values: X=500, Y=100; final values: X=200, Y=400

T1	T2	X <sub>T1</sub>	X <sub>T2</sub>	X <sub>db</sub>	Y <sub>T1</sub>	Y <sub>T2</sub>	Y <sub>db</sub>
R(X)		500		500			100
X-200		300					
W(X)	R(X)		500				
		300		300			
	X-100		400				
	W(X)		400	400			
	R(Y)				100		
					100		
R(Y)						100	
Y+200						300	
W(Y)						300	300
	Y+100					200	
	W(Y)					200	200

Conflict Serializability: Make concurrent into serial schedule with swapping orders

Example: transform a concurrent schedule to serial schedule

T1: R(A) W(A) R(B) W(B)  
T2: R(A) W(A) R(B) W(B)  
swap  
T1: R(A) W(A) R(B) W(B)  
T2: R(A) W(A) R(B) W(B)  
swap  
T1: R(A) W(A) R(B) W(B)  
T2: R(A) W(A) R(B) W(B)  
swap  
T1: R(A) W(A) R(B) W(B)  
T2: R(A) W(A) R(B) W(B)

View Serializability:  
1) Reads same version of shared object  
2) Writes same final version of shared object

T1: R(A) W(A) R(B) W(B)  
T2: R(A) W(A) R(B) W(B)

T1: R(A) W(A) R(B) W(B)  
T2: R(A) W(A) R(B) W(B)

A: T1 reads initial, T2 reads T1's write, T2 writes final  
B: T1 reads initial, T2 reads T1's write, T2 writes final

Selection: sel[expr](rel)  
-> selects the rows in some relation that satisfies expression

Projection: proj[a,b,c](rel)  
-> filters relation by its columns a, b, c

Rename: rename[schema](rel)  
-> renames the columns of rel with the columns specified in schema

sel[B=1](R) is equivalent to  
select \*  
from R  
where B = 1;  
  
Examples of selection:  
R: A B C D  
a 1 x 4  
b 2 y 5  
c 4 z 4  
d 8 x 5  
e 1 y 4  
f 2 x 5  
sel[B=1](R): A B C D  
a 1 x 4  
e 1 y 4  
sel[B=4](R): A B C D  
c 4 z 4  
sel[A=C](R): A B C D  
c 4 z 4  
sel[A=b or A=c](R): A B C D  
c 4 z 4  
e 1 y 4  
f 2 x 5

Proj[A, B, C](R) is equivalent to  
select distinct A, B, C  
from R;  
  
Examples of projection:  
R: A B C D  
a 1 x 4  
b 2 y 5  
c 4 z 4  
d 8 x 5  
e 1 y 4  
f 2 x 5  
Proj[A, B, C](R): A B C  
a 1 x 4  
b 2 y 5  
c 4 z 4  
d 8 x 5  
e 1 y 4  
f 2 x 5  
Proj[B, D](R): B D  
b 2 5  
d 8 5  
Proj[D](R): D  
4  
5  
-- Removes duplicates !!

Res = Rename[Res(b,c,d)](Project[b,c](Sel[a>5](R)) Join S)  
Tmp1 = Select[a>5](R)  
Tmp2 = Project[b,c](Tmp1)  
Tmp3 = Rename[Tmp3(cc,d)](S)  
Tmp4 = Tmp2 Join[c=cc] Tmp3  
Res = Rename[Res(b,c,d)](Tmp4)

select c.given, c.family, t.source, t.amount  
from Transactions t join Customers c on t.actor = c.id  
join Accounts a on t.source = a.id  
join Branches b on a.held\_at = b.id  
where t.tdate < '2022-05-01' and t.ttype = 'withdrawal';

Tmp1 = Sel[t.tdate < '2022-05-01' and t.ttype = 'withdrawal'] Transactions  
Tmp2 = Customers Join[Customers.id = Tmp1.actor] Tmp1  
Tmp3 = Accounts Join[Accounts.id = Tmp2.source] Tmp2  
Tmp4 = Branches Join[Branches.id = Tmp3.held\_at] Tmp3  
Result(given, family, source, amount) = Proj[given, family, source, amount] Tmp4

Union: R U S (Query1) UNION (Query2)  
Intersection: R ∩ S (Query1) INTERSECT (Query2)  
Difference: R - S (Query1) EXCEPT (Query2)

Division  
-> Here, get all A values where B has x and y.  
R: A B  
4 x  
4 y  
4 z  
5 x  
5 y  
5 z  
S: A B  
4 x  
4 y  
4 z  
5 x  
5 z  
T: B  
x  
y  
R/T: A  
4  
5  
S/T: A  
4

There is a conflict between 2 transactions if 1 & 2 apply  
1) Performs operations on same data item  
2) At least one is write (e.g. R&W, W&W, W&R)

Conflict Serializable:  
If we can swap non-conflicting operations so that result is serial schedule: 'Conflict Serializable'

View Serializable: If some schedule S and some serial schedule S'  
1. Same transaction reads same initial value of X  
2. Same transaction reads same changed version of X  
3. Same transaction writes to final value of X

e.g.  
T1: W(Y) W(X)  
T2: R(Y) W(X)  
T3: W(X)

Conflict:  
X: T1->T2, T2->T3, T1->T3  
Y: T2->T1  
Cycle In T1 and T2, not conflict serializable

View:  
Considering T2;T1;T3  
T1: W(Y), W(X)  
T2: R(Y) W(X)  
T3: W(X)

X) Final write: T3  
Y) First read: T2, Final write: T1  
Both satisfy above, therefore is 'view serializable'

<pre>having ( -- advanced(?) having sum(s.length) = ( select max(total_length) from ( select sum(s1.length) as total_length from groups g1 group by a1.id</pre>	<pre>for horse_id, horse_name in select id, name from horses where name ~* part_name order by name loop result.horse := horse_name; prize_sum := 0; races_ran := 0; select race_count into races_ran from ( select h.name, count(*) as race_count from horses h join blah where h.id = horse_id group by h.name ); for prize in (for loop in for loop) select r.prize from horses h join blah where h.id = horse_id and ru.finished = 1 loop prize_sum := prize_sum + prize; end loop; average := prize_sum / races_ran; result.average = average; return next result; end loop;</pre>	<pre>#!/usr/bin/python3 import sys import psycopg2 conn = None usage = f"Usage: {sys.argv[0]} lmao" name = sys.argv[1] conn = None  if len(sys.argv) != 2:     print(usage)     sys.exit(1)  main_query = "" select blah from blah b join what w on something order by b.lol ""  try:     conn = psycopg2.connect("dbname=bank")     cur = conn.cursor()     cur.execute(main_query, [name])     main_result = cur.fetchall() except Exception as err:     print("DB error: ", err) finally:     if conn is not None:         conn.close()</pre>	<pre>create aggregate myCount(anyelement) ( stype = int, -- the accumulator type initcond = 0, -- initial accumulator value sfunc = oneMore -- increment function );  create function oneMore(sum int, x anyelement) returns int as \$\$ begin return sum + 1; end; \$\$ language plpgsql;</pre>												
<pre>case -- case usage when instrument ilike '%guitar%' then 'guitar' when instrument in ('keyboard', 'piano') then 'piano' else instrument end  select p.id, -- string concat (p.street_no    ' '    st.name    ' '    st.stype) as street_name, su.name from properties p</pre>			<pre>create aggregate prod(numeric) ( stype = numeric, initcond = 1, sfunc = mult );  create function mult(soFar numeric, next numeric) returns numeric as \$\$ begin return soFar * next; end; \$\$ language plpgsql;</pre>												
<pre>select r.name, -- case example count(case when h.gender = 'S' then 1 end) as s_count, count(case when h.gender = 'G' then 1 end) as g_count,</pre>															
<pre>select department, string_agg(employee_name, ', ' order by employee_name) as employee_list from employees string_agg(expression, delimiter [order by expression]) group by department;</pre>															
<pre>select distinct -- distinct select b.location, c.lives_in from customers c</pre>															
<pre>select distinct -- distinct select b.location, c.lives_in from customers c</pre>		<div>return; exits function without returning further values</div>													
<pre>create or replace function q4() returns setof SongCounts as \$\$ declare result SongCounts; begin for group_id, group_name in ~ loop result."group" := group_name; for album_id in ~ loop ~ end loop; result.nshort := nshort; result.nlong := nlong; return next result;</pre>	<pre>create or replace function q3(_eventID integer) returns setof text as \$\$ declare _last integer; _tuple record; _nquitters integer := 0; begin perform id from Events where id = _eventID; if not found then return next 'No such event'; return; end if; select max(c.ordering) into _last from CheckPoints c join Events e on c.route_id = e.route_id where e.id = _eventID;  for _tuple in select person,location from farthest where event = _eventID and ordering &lt; _last loop return next _tuple.person    ' gave up at '    _tuple.location; _nquitters := _nquitters + 1; end loop;  if _nquitters = 0 then return next 'Nobody gave up'; end if; end; \$\$ language plpgsql;</pre>														
<pre>create or replace function q4() returns setof SongCounts as \$\$ declare result SongCounts; begin for group_id, group_name in ~ loop result."group" := group_name; for album_id in ~ loop ~ end loop; result.nshort := nshort; result.nlong := nlong; return next result;</pre>															
<pre>create or replace function q4() returns setof SongCounts as \$\$ declare result SongCounts; begin for group_id, group_name in ~ loop result."group" := group_name; for album_id in ~ loop ~ end loop; result.nshort := nshort; result.nlong := nlong; return next result;</pre>															
<pre>select unit_no, street_no, street, ptype into unit_num, street_num, street_id, property_type from properties where id = propID;</pre>	<pre>Meaning of ‘new’ and ‘old’ in triggers create trigger q9_2 after update of name on groups -&gt; if ‘of name’ exists, only triggered when ‘name’ column is updated.</pre>	<div>Triggers</div>													
		<pre>Triggers – what to do: 1. Split between ‘before’ and ‘after’ triggers 2. Split again by analysing the requirements  Before triggers: Raise exceptions (usually) After triggers: Update values  create trigger TotalSalary2 after update on Employee for each row execute procedure totalSalary2();  create function totalSalary2() returns trigger as \$\$ begin update Department set totSal = totSal + new.salary where Department.id = new.dept;  update Department set totSal = totSal - old.salary where Department.id = old.dept; return new; end; \$\$ language plpgsql;</pre>	<pre>create aggregate concat(text) ( stype = text, initcond = "", sfunc = join );  create function join(s1 text, s2 text) returns text as \$\$ begin if (s1 = "") then return s2; else return s1    ','    s2; end if; end; \$\$ language plpgsql;</pre>												
		<table><tr><td></td><td>Insert</td><td>Update</td><td>Delete</td></tr><tr><td>NEW</td><td>NewRow</td><td>NewRow</td><td>X</td></tr><tr><td>OLD</td><td>X</td><td>OldRow</td><td>OldRow</td></tr></table> <pre>if TG_OP = 'INSERT' then fno := new.flight_no else fno := old.flight_no end if;  if TG_OP = 'DELETE' then return old; -&gt; for ‘before delete’ else if flight.nbooked = flight.nseats then raise exception 'Booking error'; end if; -&gt; raise exception for before return new;-&gt; for ‘before insert &amp; update’  end if; P.S. returns not needed for after triggers</pre>		Insert	Update	Delete	NEW	NewRow	NewRow	X	OLD	X	OldRow	OldRow	<div>Aggregate</div>
	Insert	Update	Delete												
NEW	NewRow	NewRow	X												
OLD	X	OldRow	OldRow												
		<pre>Flights(fid, from, to, distance, departs, arrives, price) Aircraft(aid, aname, range) Certified(employee, aircraft) Employees(eid, ename, salary)  fid, aid, eid are primary keys. from,to -&gt; distance violates BCNF, we need new table</pre>	<div>SQL functions</div>												
		<div>Non-BCNF schema to BCNF schema</div>													
		<pre>new schema: Flights(fid,from,to,departs,arrives,price) Routes(from,to,distance) Aircraft(aid,aname,range) Certified(employee,aircraft) Emolovees(eid.ename.salarv)</pre>	<pre>create or replace function Q5(pattern text) returns table( rug text, size_and_stoper text, total_knots numeric(8,0) ) as \$\$ select name as rug, size    'sf '    coalesce(rug_stop::text, "") as size_and_stoper, (coalesce(knot_per_foot, 50) * coalesce(knot_per_foot, 50) * size)::numeric(8,0) as total_knots from rugs where name ~ pattern; \$\$ language sql;</pre>												
			<pre>create or replace function Q6(pattern text) returns table( province text, first integer, nrugs integer, rating numeric(3,1) ) as \$\$ select l.province, min(r.year_crafted) as first, count(r.id)::integer as nrugs, avg(r.rating)::numeric(3,1) as rating from locations l join factories f on l.id = f.located_in join crafted_by cb on f.id = cb.factory join rugs r on cb.rug = r.id where l.province ~* pattern group by l.province;</pre>												